

# Efficient Routing of Subspace Skyline Queries over Highly Distributed Data

Akrivi Vlachou, *Member, IEEE*, Christos Doulkeridis, *Member, IEEE*,  
Yannis Kotidis, and Michalis Vazirgiannis

**Abstract**—Data generation increases at highly dynamic rates, making its storage, processing, and update costs at one central location excessive. The P2P paradigm emerges as a powerful model for organizing and searching large data repositories distributed over independent sources. Advanced query operators, such as skyline queries, are necessary in order to help users handle the huge amount of available data. A skyline query retrieves the set of nondominated data points in a multidimensional data set. Skyline query processing in P2P networks poses inherent challenges and demands nontraditional techniques, due to the distribution of content and the lack of global knowledge. Relying on a superpeer architecture, we propose a threshold-based algorithm, called SKYPEER and its variants, for efficient computation of skyline points in arbitrary subspaces, while reducing both computational time and volume of transmitted data. Furthermore, we address the problem of routing skyline queries over the superpeer network and we propose an efficient routing mechanism, namely SKYPEER<sup>+</sup>, which further improves the performance by reducing the number of contacted superpeers. Finally, we provide an extensive experimental evaluation showing that our approach performs efficiently and provides a viable solution when a large degree of distribution is required.

**Index Terms**—Skyline queries, peer-to-peer systems, routing indexes.

## 1 INTRODUCTION

SKYLINE queries help users make intelligent decisions over complex data, where different and often conflicting criteria are considered. Such queries return a set of interesting data points—from the user's perspective—that are not dominated by any other point on all dimensions [1]. Skyline queries have not only been studied in centralized systems [1], but also in distributed environments such as web information systems [2] and peer-to-peer (P2P) networks [3], [4], [5], [6], [7], [8]. The P2P paradigm emerges as a powerful model for organizing and searching large data repositories distributed over independent sources [9], [10]. As an example, consider a global-scale web-based hotel reservation system, consisting of a large set of independent servers geographically dispersed around the world. Servers accept subscriptions by travel agencies in order to provide booking services over the universal hotel database, without requiring from each travel agency to register with each server. The challenge is to enable users to pose queries that capture the individual user's interests and preferences over this network of servers, and retrieve those results that match a possibly different (each time) set of user-defined criteria.

In this work, we explore how subspace skyline queries can be computed efficiently over a distributed and large-scale

web information system modeled as a P2P network. Our approach relies on a superpeer architecture, where each peer holds only a portion of the entire data set, hence horizontal data distribution is assumed. There exist several approaches for distributed skyline computation. Nevertheless, [5], [6], and [7] assume a constrained horizontal partitioning (as also mentioned in [8]) in which a structured P2P overlay network determines for each data point a server that must store it. These techniques have an expected advantage in query processing performance due to the data relocation when peer data enter the network. In this work, we make no specific assumptions and our techniques work with arbitrary horizontal partitioning of the data. In [8] and [3], the authors also employ an arbitrary horizontal partitioning, but these approaches assume the absence of an overlay network. Furthermore, in our approach, we assume a superpeer architecture that reduces the degree of decentralization compared to other P2P networks [5], but facilitates efficient query processing as also demonstrated in our experimental evaluation.

In previous work [4], SKYPEER, a distributed framework for subspace skyline processing was proposed. SKYPEER propagates the subspace skyline query to all superpeers and gathers the local skyline results sets using an efficient scheme based on threshold usage and intermediate result merging. Contacting all superpeers during query processing is practically unavoidable in the case of uniform data distribution among the superpeers, where all superpeers may store part of the skyline result set. But in many real-life applications, data are clustered and then it is imperative to avoid contacting all superpeers. In this paper, we address the problem of routing the skyline queries over a superpeer network, aiming to reduce the number of contacted superpeers. Each superpeer clusters its local data, in order to provide a summary description and to make its contents searchable by other superpeers. The cluster descriptions are

• A. Vlachou and C. Doulkeridis are with the Department of Computer and Information Science (IDI), Norwegian University of Science and Technology (NTNU), Sem Saelandsv. 7-9, N-7491 Trondheim, Norway. E-mail: {vlachou, cdoulk}@idi.ntnu.no.

• Y. Kotidis and M. Vazirgiannis are with the Department of Informatics, Athens University of Economics and Business, 76 Patission Street, GR 10434 Athens, Greece. E-mail: {kotidis, mvazirgi}@aueb.gr.

Manuscript received 21 May 2008; revised 15 Dec. 2008; accepted 3 Sept. 2009; published online 20 Nov. 2009.

Recommended for acceptance by G. Kollios.

For information on obtaining reprints of this article, please send e-mail to: tkde@computer.org, and reference IEEECS Log Number TKDE-2008-05-0277. Digital Object Identifier no. 10.1109/TKDE.2009.204.

published to the other superpeers and routing indexes are built. Therefore, we propose an appropriate indexing and routing mechanism for subspace skyline queries over a superpeer network. The main contributions of our work are:

- We explore the implications of processing and routing subspace skyline queries in large-scale P2P networks.
- We present SKYPEER which utilizes a thresholding scheme in order to facilitate pruning of dominated data across the peers. We explore different strategies for 1) threshold propagation and 2) result merging over the P2P network.
- We introduce SKYPEER<sup>+</sup>, an efficient routing algorithm, which improves performance in the case of clustered distributions of data among superpeers. By using the routing indexes, SKYPEER<sup>+</sup> further improves the thresholding scheme and reduces drastically the transferred data.
- We present a set of novel techniques for efficient management of cluster information. We propose the iSUBSKY mapping for local indexing of clusters, threshold refinement for individual clusters and cluster-driven query routing.
- We provide an extensive experimental evaluation of our algorithms and show their comparative benefits over a baseline skyline computation algorithm adapted to work in P2P networks. In addition, we provide an experimental comparison with existing algorithms SSP [5] and FDS [8] for various setups.

The rest of this paper is organized as follows: Section 2 provides an overview of related work. In Section 3, we present the necessary definitions and we introduce the extended skyline. In Section 4, we give an overview of the P2P system model and sketch our approach. In Section 5, we describe the preprocessing phase. Then, in Section 6, we present the basic SKYPEER algorithm, and in Section 7, we introduce the efficient routing algorithm SKYPEER<sup>+</sup>. In Section 8, we provide a comparative performance analysis and we discuss issues related to maintenance. Our experimental results are presented in Section 9. Finally, in Section 10, we conclude the paper.

## 2 RELATED WORK

Computing skyline was first investigated in computational geometry [11]. Börzsönyi et al. [1] first investigate the skyline computation problem in the context of databases. Thereafter, several techniques are proposed in the relevant research literature for the efficient skyline computation in a centralized setting [12], [13], [14], [15], [16]. Motivated by the different preferences of users, recent papers focus on algorithms to support *subspace skyline* retrieval. Pei et al. [17] and Yuan et al. [18] independently propose the SKYCUBE, which consists of the skylines in all possible subspaces. In [19], [20], SUBSKY algorithm is presented, which transforms the multidimensional data to one-dimensional values, and then indexes the data set with a B-tree. Similarly, in [16], the idea of limiting the amount of data to be read by exploiting the value of a monotone function was studied. However, [16] does not study subspace skyline queries.

The skyline operator has been also studied in decentralized and distributed environments. In [2], an algorithm for

distributed Web-accessible sources was proposed. Unfortunately, their assumptions are hardly applicable to large-scale P2P systems. Hose et al. [21] focus on P2P skyline computation, while providing probabilistic guarantees for the result's correctness. In comparison, our algorithms provably return *exact* answers to arbitrary subspace skyline computations. In [22], Huang et al., assume a setting with mobile devices communicating via an ad hoc network, and study skyline queries that involve spatial constraints. The problem of parallelizing progressive constrained skyline queries [14] in a shared-nothing architecture is addressed in [6]. Although unconstrained skyline queries could probably be supported, this approach suffers from poor load balancing and additional mechanisms, such as replication, need to be used.

An approach for P2P skyline query processing using a tree-structured overlay is presented in [5] and later extended in [23]. In this approach, due to the employed space partitioning scheme, the problem of load balancing is quite important. In [7], an approach for P2P skyline retrieval is described that uses Chord as the underlying infrastructure and minimizes the bandwidth consumption as well as the number of visited nodes. In [24], an approach for continuous subspace skylines in a distributed setting that is based on bitmap representation is presented. In recent work [3], the *PaDSkyline* algorithm is proposed for constrained skyline query processing in distributed environments. However, in contrast to our approach, the authors do not address the issue of communication among sites (peers), as they assume that no overlay exists. In [8], the FDS algorithm is also proposed for skyline query processing in distributed environments where no overlay exists. The aim is to minimize the number of transferred objects; however, this may result in several round-trips until the correct result is computed, thereby, incurring a high total response time.

Routing indexes for efficient search in P2P systems have been previously proposed in the research literature [25], [26]. In this paper, we capitalize on such techniques, in order to build an efficient routing mechanism at superpeer level.

## 3 PRELIMINARIES AND DEFINITIONS

We assume a P2P network of  $N_p$  peers, where some peers have special roles, due to enhanced features, such as availability, storage capability, and bandwidth capacity. These peers are called superpeers  $SP_i$  ( $i = 1..N_{sp}$ ), and they constitute only a small fraction of the peers in the network, i.e.,  $N_{sp} \ll N_p$ . Peers that join the network directly connect to one of the superpeers. Each superpeer maintains links to simple peers, based on the value of its degree parameter  $DEG_p$ . In addition, a superpeer is connected to a limited set of at most  $DEG_{sp}$  other superpeers ( $DEG_{sp} < DEG_p$ ).

Each peer  $P_i$  holds  $n_i$   $d$ -dimensional points, denoted as a set  $S_i$  ( $i = 1..N_p$ ). Obviously, the size of the complete set of points is  $n = \sum_{i=1}^{N_p} n_i$  and the data set  $S$  is the union of all peers' data sets  $S_i$ :  $S = \cup S_i$ . Given a space  $D$  defined by a set of  $d$  dimensions  $\{d_1, d_2, \dots, d_d\}$  and a data set  $S$  on  $D$ , a point  $p \in S$  can be represented as  $p = \{p[1], p[2], \dots, p[d]\}$  where  $p[i]$ , is a value on dimension  $d_i$ .

**Definition 3.1 (Skyline).** A point  $p \in S$  is said to dominate another point  $q \in S$ , denoted as  $p \prec q$ , if 1) on every dimension




Fig. 1. Skyline example.

$d_i \in D$ ,  $p[i] \leq q[i]$ , and 2) on at least one dimension  $d_j \in D$ ,  $p[j] < q[j]$ . The skyline is a set of points  $SKY \subseteq S$  which are not dominated by any other point. The points in  $SKY$  are called skyline points.

Without loss of generality, we assume that skylines are computed with respect to min conditions on all dimensions and that all values are non-negative. The notion of skyline can be extended to subspaces. Each nonempty subset  $U$  of  $D$  ( $U \subseteq D$ ) is referred to as a *subspace* of  $D$ . The data space  $D$  is also referred to as full space of the data set  $S$ .

**Definition 3.2 (Subspace Skyline).** A point  $p \in S$  is said to dominate another point  $q \in S$  on subspace  $U \subseteq D$ , denoted as  $p \prec_U q$ , if 1) on every dimension  $d_i \in U$ ,  $p[i] \leq q[i]$ , and 2) on at least one dimension  $d_j \in U$ ,  $p[j] < q[j]$ . The skyline of a subspace  $U \subseteq D$  is a set of points  $SKY_U \subseteq S$  which are not dominated by any other point on subspace  $U$ . The points in  $SKY_U$  are called skyline points on subspace  $U$ .

Consider, for example, the data set depicted in Fig. 1. The skyline points are  $SKY = \{a, i, k\}$ , while for the (nonempty) subspace  $U = \{y\}$  the skyline points on  $U$  are  $SKY_U = \{k, d\}$ . Notice that the point  $d$  is a skyline point on the subspace  $\{y\}$  but it is dominated by the point  $k$  in the space  $\{x, y\}$ . As shown in [18], [4], the skyline set of the full space does not contain all the skyline points for any subspace. A skyline point  $q$  in  $SKY_U$  is either a skyline point in  $SKY_V$  (assuming  $U \subset V$ ) or there is another data point  $p$ , such that  $p[i] = q[i]$  ( $\forall d_i \in U$ ), that dominates  $q$  on the dimension set  $V - U$ . Thus, a superset of the union of all subspace skylines is the skyline set of the full space enriched with all points  $p$  for which there exists a point  $q \in SKY_D$  such that  $q[i] = p[i]$  for at least one dimension  $d_i$ . Consider, for example, Fig. 1 where  $e$  and  $k$  have the same  $x$ -value but  $k$  is a subspace skyline point in contrast to point  $e$  which is not a skyline point in any subspace.

To reduce the computation overhead, we define a subset of the aforementioned superset that is able to answer any subspace skyline query. Therefore, we adjust the dominance definition to compute all necessary values coinstantaneously during the skyline calculation. We define [4] the *extended skyline* (*ext-skyline*) based on the *extended domination* (*ext-domination*) definition.

**Definition 3.3 (Extended Skyline).** For any dimension set  $U$ , where  $U \subseteq D$ ,  $p$  *ext-dominates*  $q$  if on each dimension  $d_i \in U$ ,  $p[i] < q[i]$ . The *ext-skyline* (*ext-SKY<sub>U</sub>*) is set of all points that are not *ext-dominated* by any other.

The following lemmas show that *ext-SKY<sub>D</sub>* is sufficient to answer any subspace skyline query correctly:

**Lemma 3.4.** Every point that belongs to the skyline of  $U$  belongs also to the *ext-skyline* of  $U$ , i.e.,  $SKY_U \subseteq \text{ext-SKY}_U$ .

**Lemma 3.5.** Every point that belongs to the skyline of a subspace  $V \subseteq U$  belongs to the *ext-skyline* of  $U$ , i.e.,  $SKY_V \subseteq \text{ext-SKY}_U$ ,  $V \subseteq U$ .

For example, in Fig. 1, points  $b$ ,  $m$ , and  $d$  belong to the *ext-skyline*, which is not the case with point  $e$ , since  $e$  is globally dominated by  $i$ , which in turn does not have any value equal to the attribute values of  $e$ . Notice that neither  $e$  nor  $m$  belong to any subspace skyline, in contrast to  $b$  and  $d$ .

## 4 SYSTEM OVERVIEW

In the following, we first describe a baseline approach, henceforth referred to as *naive*. Then, we describe our thresholding approach, which is more efficient than the naive approach. Finally, we outline an extension of our algorithm with enhanced routing capabilities and particularly suitable for clustered data distributions.

### 4.1 A Baseline Algorithm

In a preprocessing phase (Section 5.2), each peer  $P_i$  computes the local *ext-skyline* of its data set  $S_i$ . Each superpeer  $SP_i$ , collects the *ext-skylines* of its associated peers and merges them by discarding all *ext-dominated* points. Thereafter, given a subspace skyline query, each superpeer is able to answer the query based on its peers' data without actually contacting any peer. The querying peer  $P_{init}$  is usually a simple peer that submits a query to the system. However, in the rest of this paper, we use  $P_{init}$  to refer to the superpeer responsible for the simple peer. In order to answer a subspace skyline query over the entire superpeer network, the initiator superpeer  $P_{init}$  has to contact all other superpeers through the neighboring superpeers. More detailed, each superpeer that receives a subspace skyline query, individually processes the request based on the locally stored *ext-skylines* and routes the results back to the query initiator through its neighbors. The query initiator collects the results from all superpeers and merges them by discarding dominated points. This approach is referred to as *naive*. In what follows, we describe our approach, which aims at improving the performance of the naive approach, by drastically reducing communication costs and the overall processing time.

### 4.2 Threshold Usage for Skyline Queries

Subspace skyline points computed at any superpeer may dominate—and thus prune—points of the current superpeer. Therefore, in our algorithm, a threshold value is defined based on already computed subspace skyline points, and the threshold is attached to the query before it is propagated in the network. The threshold can be updated at any superpeer in the network during query processing. In our work, we explore different strategies for threshold propagation and result merging through the P2P network (Section 6). In order to support threshold-based query processing, data is transformed into one-dimensional values in a preprocessing phase (Section 5).

### 4.3 Routing of Skyline Queries

Contacting all superpeers during query processing is acceptable for the case of uniform data distribution among superpeers, since all superpeers probably store some skyline points and have to be contacted, in order to retrieve the exact skyline result set. But in the case of clustered data distribution, it is desirable to avoid contacting superpeers that do not contribute to the result set.

Instead of flooding the network, we propose a routing mechanism to contact only those superpeers that may contribute to the overall subspace skyline result set, while reducing the number of contacted superpeers and transferred data. More detailed, in the preprocessing phase, each superpeer  $SP_A$  additionally applies a clustering algorithm on its ext-skyline points. The cluster descriptions are broadcast over the superpeer network. Each superpeer collects the cluster information of all superpeers and builds routing indexes based on them (Section 7.3). During query processing, the routing indexes are used to define which superpeers may contribute to the overall subspace skyline set (Section 7.4), thus, avoiding flooding.

## 5 MAPPING AND PREPROCESSING

At superpeer level there exists a preprocessing phase that allows subsequent query processing over the entire network of peers. In order to enable the threshold-based skyline computation, the multidimensional data are transformed into one-dimensional values. This mapping is the topic of the next section. Thereafter, we discuss the preprocessing phase in detail.

### 5.1 Mapping

Recall that each peer maintains a portion of the data set. A query refers to a nonempty subset  $U$  of  $D$ . Inspired by [19], [16], each  $d$ -dimensional point  $p$  is transformed to a one-dimensional value  $f(p)$  based on the formula:

$$f(p) = \min_{i=1}^d (p[i]). \quad (1)$$

Let  $dist_U(p)$  denote the  $L_\infty$ -distance of point  $p$  from the origin based on the dimension set  $U$ , i.e.,  $dist_U(p) = \max_{i \in U} (p[i])$ .

**Observation 5.1.** Let  $p_{sky}$  be a skyline point in a subspace  $U$ . A point  $p$  for which the following inequality holds cannot be a skyline point in subspace  $U$ .

$$f(p) > dist_U(p_{sky}). \quad (2)$$

Algorithm 1 uses the proposed mapping for efficient local subspace skyline computation at a peer. Data points are accessed in an ascending order of their  $f(p)$  values. Note that the  $f(p)$  value is computed once on  $D$  independently from the queried subspace  $U$ . In contrast,  $dist_U(p)$  refers to  $U$  and it is calculated during the skyline computation. The nondominated points among the data already examined are added into the current subspace skyline set  $SKY_U$ . The algorithm uses as *threshold* the minimum value of the  $dist_U(p)$  of all points in  $SKY_U$ . Based on Observation 5.1, the algorithm terminates when *threshold* is smaller than the  $f(p)$  value of the next point  $p$ .




Fig. 2. Mapping example.

#### Algorithm 1. Subspace skyline computation

```

1: Input:  $U$ : query dimensions
    $L$ : sorted list of data points
2:  $SKY_U \leftarrow \emptyset$ 
3:  $threshold \leftarrow MAX\_INT$ 
4:  $p \leftarrow L.pop()$ 
5: while ( $f(p) < threshold$ ) do
6:   if ( $\nexists q \in SKY_U : q \prec_U p$ ) then
7:      $SKY_U \leftarrow SKY_U - \{q\}, p \prec_U q$ 
8:      $SKY_U \leftarrow SKY_U \cup \{p\}$ 
9:      $threshold \leftarrow \min_{p_i \in SKY_U} (dist_U(p_i))$ 
10:  end if
11:   $p \leftarrow L.pop()$ 
12: end while
13: return  $SKY_U$ 

```

Fig. 2 depicts a mapping example for a two-dimensional data set. For sake of simplicity, in our examples, we assume that the query space and the data space are identical. In Fig. 2, the dashed lines correspond to the points with  $f(p)$  values of 1, 2, and 3, respectively. Actually, we examine the data space in a way that is equivalent to the dashed line being shifted from the origin toward the right corner of the data space. On the other hand, the dotted line shows the points that cause a threshold value ( $dist_U(p_{sky})$ ) of 3 and 4, respectively. Notice that if there exists a point that sets the threshold equal to 3 (lies on the corresponding dotted line), this point dominates any point on the dashed line with  $f(p)$  equal to 3 independently of the exact location of the points (Observation 5.1). Therefore, the threshold-based algorithm manages to return the correct skyline set.

### 5.2 Preprocessing Phase

First, each peer  $P_i$  computes the local ext-skyline of its data set  $S_i$  and sends it to the associated superpeer. Then, the superpeer gathers the local ext-skylines of the individual peers and merges them by pruning out those points of a peer  $P_i$  that are ext-dominated by points of another peer  $P_j$ , resulting in one set that constitutes the ext-skyline of space  $D$  with respect to the data set on the superpeer and its associated peers. Based on Lemma 3.4, the local ext-skyline is sufficient for a superpeer to determine if any of its peers  $P_i$  contributes to the results of any subspace skyline query.

We illustrate the details of peer preprocessing by means of an example. In Fig. 3, three peers ( $P_A, P_B, P_C$ ) are assigned to superpeer  $SP_A$  and their local data sets are shown. The dimensionality of the data set is 4. Each peer computes its local ext-skyline in the original space. The points added to the result set due to the ext-skyline definition are gray




Fig. 3. Peer preprocessing example.

shaded. For instance, four of the five points of  $P_A$  are skyline points, while  $A3$  is included as an ext-skyline point.

## 6 THE BASIC SKYPEER ALGORITHM

Our algorithm utilizes a thresholding scheme, in order to facilitate pruning of dominated data across the peers. A querying superpeer hands on the query to its neighboring superpeers along the superpeer backbone, which in turn forward the query to their adjacent superpeers. The superpeers execute the query over locally stored ext-skylines and retrieve local results, which are sent back through the query routing path. In the sequel, we present all relevant steps in detail.

### 6.1 Threshold-Based Skyline Processing

Let  $t$  be the threshold value at the end of the local skyline computation (i.e., Algorithm 1) at the initiator. Based on Observation 5.1, the threshold value indicates that there is a local skyline point  $p$  that dominates all points with  $f(p)$  values larger than  $t$ . At the end of the local skyline computation,  $t$  corresponds to the point with the minimum threshold value, i.e., the highest pruning capability. Since data are horizontally partitioned over the superpeers, the local skyline point  $p$  dominates all points of any superpeer with  $f(p)$  larger than  $t$ . Therefore,  $t$  is attached to the query and propagated to the superpeers. The threshold value is used as an initial threshold for the local subspace skyline computation to further reduce the computation and communication cost.

An alternative threshold propagation strategy is to compute and refine the threshold on each superpeer that is processing the query, instead of forwarding  $P_{init}$ 's fixed threshold value. Intuitively, by progressively lowering the threshold value, the pruning capability of the query increases at each forwarding step. However, this approach requires that the query is propagated after the superpeer has finished the local skyline computation.

### 6.2 Merging Phase

During query processing, the local subspace skyline sets have to be merged into one overall result set. We explore two different strategies for the merging phase.

The simplest strategy is that the initiator superpeer  $P_{init}$  collects the local subspace skyline result sets of all superpeers

TABLE 1  
Variants of the Basic SKYPEER

Fixed Threshold Fixed Merging (FTFM)
Fixed Threshold Progressive Merging (FTPM)
Refined Threshold Fixed Merging (RTFM)
Refined Threshold Progressive Merging (RTPM)

through its neighbors and merges the local result sets of the individual superpeers to one global result set. Even though the use of the threshold reduces the amount of data transmitted to the query initiator  $P_{init}$ , there is still a chance that a local result may contain points that do not belong to the overall skyline. Therefore, a second strategy for the merging phase is to merge progressively the local skyline sets during query evaluation. Instead of forwarding all results back to  $P_{init}$ , each superpeer merges the results of its neighbors, and forwards the merged result back to  $P_{init}$ . The benefit is twofold. The transferred data are reduced and a time-consuming centralized merging phase is avoided. On the other hand, the computation cost for each superpeer increases due to the additional merging phase. However, we note that the input and, therefore, the cost of each subsequent merging phase is much smaller.

### 6.3 SKYPEER

SKYPEER propagates the subspace skyline query to all superpeers and gathers the local subspace skyline result sets using an efficient thresholding scheme. Thereafter, these local subspace skyline results have to be merged into a global result set, for example, by the initiator superpeer. In Table 1, we present our variants of the basic SKYPEER based on two different optimization criteria for distributed subspace skyline computation:

1. Threshold propagation: a) *Fixed Threshold*:  $P_{init}$  calculates its threshold  $t$  for  $q(U, t)$  and forwards the threshold value to all superpeers. b) *Refined Threshold*:  $P_{init}$  calculates and sends its threshold to its neighboring superpeers, which do not forward it immediately to other superpeers, but rather they first compute the subspace skyline, calculate the new threshold  $t'$ , and then forward  $q(U, t')$ .
2. Merging strategy: a) *Fixed Merging at  $P_{init}$* : In this approach, all superpeers forward their computed subspace skyline back to  $P_{init}$ , and  $P_{init}$  is responsible for merging the results and computing the resulting subspace skyline for  $q(U)$ . b) *Progressive Merging*: Each superpeer merges the results it receives with its locally computed subspace skyline, before sending the results back to the superpeer from which it originally received the query.

We now introduce the SKYPEER algorithm (see Algorithm 2). A subspace skyline query  $q(U)$  is posed by the initiator superpeer  $P_{init}$ . The initiator superpeer first computes the skyline on its local ext-skyline resulting in a threshold value  $t$ , which based on Observation 5.1 can be used to prune out points that cannot belong to the result of the skyline query. The threshold value is attached to the query  $q(U)$  that becomes  $q(U, t)$  before it is forwarded to  $P_{init}$ 's neighbors at superpeer level. Each

superpeer  $SP_i$  receiving the query, forwards it to its neighboring superpeers and executes a local *threshold-based* subspace skyline computation on its local ext-skyline points. If the RT\*M variants are employed, then before forwarding the query,  $SP_i$  computes the skyline on its local ext-skyline which results in a refined threshold value  $t'$  (or in worst case  $t' = t$ ) that is attached to the query before it is sent to the neighboring superpeers. The results are routed back to  $P_{init}$ . If one of the \*TPM variants is employed,  $SP_i$  first merges the results it receives, before forwarding them to the superpeer from which it received the query (querying superpeer).

**Algorithm 2.** SKYPEER on superpeer  $SP_i$

```

1: Input: mode (FTFM, FTPM, RTFM, RTPM)
   Query ( $q(U, t)$ )
   Querying superpeer ( $SP_q$ )
2:  $L_N \leftarrow$  list of all neighbors except  $SP_q$ 
3: if (mode  $\in$  {RTFM, RTPM}) or ( $SP_i = P_{init}$ ) then
4:    $SKY_{U_0} \leftarrow$  compute local skyline
5:    $t \leftarrow refineThreshold(t)$ 
6: end if
7: for  $n_i \in L_N$  do
8:   send( $n_i, q(U, t)$ )
9: end for
10: if mode  $\in$  {FTFM, FTPM} then
11:    $SKY_{U_0} \leftarrow$  compute local skyline
12: end if
13: for  $j = 1$  to  $|L_N| - 1$  do
14:   receive  $SKY_{U_j}$ 
15:   if mode  $\in$  {FTPM, RTPM} then
16:      $SKY_{U_0} \leftarrow mergeResults(SKY_{U_j}, SKY_{U_0})$ 
17:   else
18:     send( $SP_q, SKY_{U_j}$ )
19:   end if
20: end for
21: send( $SP_q, SKY_{U_0}$ )

```

## 7 THE SKYPEER<sup>+</sup> ALGORITHM

In this section, we describe SKYPEER<sup>+</sup>, an algorithm for the case of nonuniform data distribution on peers, that aims at efficient and deliberate routing of skyline queries. In order to achieve this goal, a mechanism for pruning superpeers that cannot contribute to the result is necessary. In our approach, superpeers exchange compact data descriptions, thereby, enabling pruning of dominated superpeers at query time. The exchanged data descriptions enable the construction of routing indexes for intentional query forwarding to specific superpeers.

Each superpeer clusters the extended skyline points using a standard clustering algorithm. During the preprocessing phase, each superpeer publishes only the cluster descriptions  $C_i$  to all other superpeers, as a summarization of its data, while the extended skyline is stored by the superpeer itself. The remaining challenge is to answer skyline queries over the entire superpeer network. Instead of flooding queries at superpeer level, we build routing indexes based on the cluster descriptions that enable selective query routing only to superpeers that may actually be responsible for peers with relevant results.

In Section 7.1, we introduce the one-dimensional mapping used by SKYPEER<sup>+</sup>, while in Section 7.2, we discuss how local query processing is performed on a superpeer. In Section 7.3, we describe the construction of the routing indexes. Then, in Section 7.4, we present the cluster dominance relationships and we introduce the skyline routing algorithm.

### 7.1 Mapping and Representation

Recall that each superpeer stores the extended skylines of its associated peers, henceforth, also called data points of the superpeer.

For nonuniform data sets, we follow the same intuition of [19] where the data set is partitioned in  $N_C$  clusters. Each superpeer first clusters the extended skyline points using a standard clustering algorithm (like K-Means) or using an application specific clustering method [19]. Thus, each data object is assigned to the nearest cluster based on the distance to the cluster's centroid. The superpeer determines for each cluster  $C_i$  the minimum bounding rectangle (MBR), which is represented by two reference points  $l_i$  and  $u_i$ . Point  $l_i$  is defined as  $l_i[j] = \min_{p \in C_i} (p[j])$  and dominates all points in  $C_i$ , while point  $u_i$  is defined as  $u_i[j] = \max_{p \in C_i} (p[j])$  and is dominated by all points in  $C_i$ . For each data point  $p \in C_i$ , a one-dimensional mapping is applied according to the distance of  $p$  to  $l_i$ :

$$f(p, C_i) = \min_{j=1}^d (p[j] - l_i[j]). \quad (3)$$

The one-dimensional mapping combined with the clustering information allows us to determine which data points cannot belong to the subspace skyline set and can therefore safely be pruned. Let  $dist_U(p, C_i)$  denote the  $L_\infty$ -distance of point  $p$  that belongs to the cluster  $C_i$  from the lower corner  $l_i$  of the  $MBR_i$  based on the dimension set  $U$ , i.e.,

$$dist_U(p, C_i) = \max_{j \in U} (p[j] - l_i[j]). \quad (4)$$

**Observation 7.1.** Let  $p_{sky}$  be a skyline point in a subspace  $U$ . A point  $p \in C_i$  for which the following inequality holds cannot be a skyline point in subspace  $U$ .

$$f(p, C_i) > dist_U(p_{sky}, C_i). \quad (5)$$

Based on Observation 7.1, already examined points define a threshold  $t(C_i)$  for each cluster  $C_i$ , which indicates the region within the cluster that is pruned. Notice that even a point  $p$  that does not belong to cluster  $C_i$  can dominate points that belong to  $C_i$ , as long as  $p$  dominates  $u_i$ , and thus  $p$  can refine the threshold  $t(C_i)$ .

Consider, for example, Fig. 4. Point  $p$  for which  $f(p, C_1) = 1$  sets the threshold  $t(C_1)$  based on (4) equal to 2. Therefore, the pruned area of the cluster  $C_1$  is the dark area within  $C_1$ . The dominance area of  $p$  partially covers cluster  $C_2$ , since  $p$  dominates  $u_2$  and the threshold value  $t(C_2)$  may be refined by  $p$ . Hence, the threshold  $t(C_2)$  is set to 3, i.e.,  $dist_U(p, C_2)$ , indicating that point  $p$  prunes the dashed area in cluster  $C_2$ . In this way, we can prune points in cluster  $C_2$  using the threshold, before we even examine any point of cluster  $C_2$ .

### 7.2 Indexing and Local Query Processing

In the following, we describe the indexing method, namely iSUBSKY (Section 7.2.1), that is employed by the superpeer to




Fig. 4. Threshold example.

index its data, in order to answer efficiently subspace skyline queries (Section 7.2.2) during local query processing. The iSUBSKY method indexes the extended skyline points at each superpeer, since these points are sufficient for a superpeer to answer a subspace skyline query locally. During subspace skyline query processing, iSUBSKY enables early pruning of clusters that do not contribute to the final result set, which reduces the processing cost significantly.

### 7.2.1 Indexing

Our approach uses a thresholding scheme (Observation 7.1) that relies on a one-dimensional mapping, in order to efficiently prune points and clusters that do not belong to the result set. The one-dimensional values  $f(p, C_i)$  refer to the full space (original data space) and do not depend on the queried subspace; therefore, they can be computed in a preprocessing phase and stored on disk. To maintain efficiently the one-dimensional values, the values have to be stored to a one-dimensional index, such as  $B^+$ -tree. In addition, we need a representation of the values where each cluster corresponds to a separate interval in the  $B^+$ -tree.

Inspired by [27], [9], [19], we propose iSUBSKY, which is an indexing method suitable for subspace skyline query processing. Given a set of  $N_C$  clusters and a constant  $c$ , each data object  $p$  that belongs to a cluster  $C_i$  is assigned a one-dimensional iSUBSKY value according to its  $f(p, C_i)$  value:

$$iSUBSKY(p) = i * c + f(p, C_i). \quad (6)$$

Using a sufficiently high  $c$  value, all objects in the  $i$ th cluster  $C_i$  are mapped to the interval  $[i * c, i * c + f(u_i, C_i)]$ , which is nonoverlapping with other cluster intervals. Fig. 5 depicts an example of the iSUBSKY mapping. Notice that points are mapped into one-dimensional values, whereas clusters correspond to intervals. The actual data points can now be efficiently stored in a  $B^+$ -tree using the iSUBSKY values as keys. Additionally, the clusters  $C_i = \{l_i, u_i\}$  are kept in a main memory list  $C$ . In this way, we can process a subspace skyline query by examining only some specific intervals of the  $B^+$ -tree, as described in the following paragraphs.

### 7.2.2 Subspace Skyline Query Processing

During local query processing on a superpeer, the cluster information is accessed from the main memory list  $C$  and the points that belong to each cluster are retrieved using the index. For each cluster  $C_i$  accessed at query time, we have to examine at most the interval  $[i * c, i * c + f(u_i, C_i)]$ . Similar to Algorithm 1, the points for each cluster  $C_i$  are examined in an ascending order of  $f(p, C_i)$  values. We also




Fig. 5. Example of iSUBSKY.

keep a threshold value  $t(C_i)$  for each cluster  $C_i$ . Notice that most probably we do not have to retrieve all data points in the searched interval, as the processing may stop earlier based on the threshold condition (Observation 7.1).

Algorithm 3 describes the pseudocode in detail. For a subspace skyline query  $q(U)$ , the cluster descriptions  $C_i$  are kept in a main memory list  $C$  of clusters sorted in ascending order based on the values  $\min_{j \in U} (l_i[j])$  (Line 3). In each iteration, we examine the next cluster  $C_m$  (Line 7) and we retrieve the point  $p$  that belongs to  $C_m$  with the minimum iSUBSKY ( $f(p, C_m)$ ) value (Lines 9 and 10). We repeatedly retrieve the next point based on the iSUBSKY value, until the threshold condition holds (Line 11), since we can safely ignore the remaining points. Each time a new point  $p$  is retrieved, we examine if this point is dominated by any point in  $SKY_U$  or by any cluster  $C_i \in C$  (Line 12).<sup>1</sup> If this is not the case, we remove all points  $q$  retrieved so far that are dominated by  $p$  (Line 13), as well as all clusters  $C_i$  that are dominated by  $p$  (Line 14). Then,  $p$  is added to  $SKY_U$  as a candidate skyline (Line 15). We also refine the current threshold (Lines 16-18) of all clusters  $C_i$  for which  $p$  dominates  $u_i$ . The threshold  $t(C_i)$  is set as the minimum value of the previous threshold  $t(C_i)$  and the new distance  $dist_U(p, C_i)$ . In the next iteration, if  $f(p, C_m) > t(C_m)$  we discard  $C_m$  without retrieving the remaining points of the cluster from the  $B^+$ -tree (Line 11). Our algorithm returns the subspace skyline set  $SKY_U$  (Line 24), after examining all clusters that are not dominated by any skyline point retrieved so far.

#### Algorithm 3. Superpeer subspace skyline computation

- 1: **Input:**  $U$ : query dimensions  
 $C$ : list of clusters  $\{C_i\}$
- 2: **Output:**  $SKY_U$
- 3: Sort  $C$  in ascending order according to the value  $\min_{j \in U} (l_i[j])$  for each  $C_i$
- 4:  $SKY_U \leftarrow \emptyset$
- 5:  $t(C_i) \leftarrow \text{MAX\_INT}, \forall C_i \in C$
- 6: **while** ( $C \neq \emptyset$ ) **do**
- 7:  $C_m \leftarrow C.\text{pop}()$
- 8:  $C \leftarrow C - \{C_i\}, C_m \prec_U C_i$
- 9:  $cursor \leftarrow \text{search}([m * c, m * c + f(u_m, C_m)])$
- 10:  $p \leftarrow cursor.\text{pop}()$

1. A cluster  $C_i(l_i, u_i)$  dominates a point  $p$  in  $U$ , if  $u_i \prec_U p$ .




Fig. 6. Routing indexes construction.

```

11: while ( $f(p, C_m) < t(C_m)$ ) and ( $p \neq \text{null}$ ) do
12:   if ( $\nexists q \in SKY_U : q \prec_U p$ ) and
      ( $\nexists C_i \in C : C_i \prec_U p$ ) then
13:      $SKY_U \leftarrow SKY_U - \{q\}, p \prec_U q$ 
14:      $C \leftarrow C - \{C_i\}, p \prec_U C_i$ 
15:      $SKY_U \leftarrow SKY_U \cup \{p\}$ 
16:     for ( $\forall C_i \in C : p \prec_U u_i$ ) do
17:        $t \leftarrow \max_{j \in U} (p[j] - l_i[j])$ 
18:        $t(C_i) \leftarrow \min(t(C_i), t)$ 
19:     end for
20:   end if
21:    $p \leftarrow \text{cursor.pop}()$ 
22: end while
23: end while
24: return  $SKY_U$ 

```

Instead of iSUBSKY, the method proposed in [19] could be used that stores tuples of cluster identifiers and  $f(p)$  values in a B<sup>+</sup>-tree. This method applies a heuristic to decide which cluster to examine in each step and to reduce the threshold of all clusters, while exploring all clusters in parallel by using multiple iterators. Our proposed method, examines the clusters in a serialized way, so that the threshold is reduced only by already examined clusters. The main advantage of iSUBSKY is that it allows us to discard entire clusters without accessing any of their data points.

### 7.3 Routing Indexes Construction

During the preprocessing phase (Section 5), each superpeer additionally broadcasts its cluster descriptions  $C_i(l_i, u_i)$  to the network, in order to build routing indexes at superpeer level. Notice that this is a tolerable cost, since it is a one-time cost (when no updates occur) and it does not saturate the network.

After the exchange, each superpeer holds a list of clusters received through each neighbor. In order to efficiently maintain and process the collected cluster information, we use the iSUBSKY method for indexing the cluster information. Each superpeer applies a clustering algorithm<sup>2</sup> on all received clusters, resulting in  $N_{RC}$  routing clusters  $RC_i(rl_i, ru_i)$ , which are represented as MBRs defined by the left lower corner  $rl_i$  and the right upper




Fig. 7. Cycles example.

corner  $ru_i$  of the MBR. Fig. 6 depicts an example where two routing clusters  $RC_1, RC_2$  are created that summarize four clusters  $C_1-C_4$ . The routing clusters  $RC_i$  are kept in main memory. For each cluster  $C_j$  that belongs to a routing cluster  $RC_i$ , we index its lower corner ( $l_j$ ) using the iSUBSKY value  $f(l_j, RC_i)$ , and we store for each cluster a triple  $\{f(l_j, RC_i), l_j, u_j\}$ . The triples are stored in a B<sup>+</sup>-tree using the iSUBSKY values as keys. By accessing the clusters  $C_j$  in ascending order of the  $f(l_j, RC_i)$  value and setting the threshold based on the maximum distance to the upper corner  $u_j$ , it is easy to extend Algorithm 3 for subspace skyline computation, where instead of data points, the dominance relationships between MBRs are considered. In the next section, we will further discuss, how efficient query routing can be established, while computing additionally a threshold for each nondominated cluster. Using the routing clusters, a superpeer can determine to which neighbors a query  $q(U)$  should be forwarded, and thus prune network paths represented by dominated clusters.

Query routing through a superpeer network that contains cycles leads to redundant messages. Therefore, to improve the performance of our routing indexes, each superpeer keeps a list  $L_i$  for each superpeer  $SP_i$ , which contains superpeer identifiers. A superpeer identifier  $SP_i \in L_s$  at superpeer  $SP_j$  can be interpreted as: clusters of  $SP_s$  are reached by  $SP_i$  through a path that contains  $SP_j$ . Thus,  $SP_j$  forwards queries that refer to clusters of  $SP_s$ , only if the query initiator belongs to  $L_s$ .

During the construction of routing indexes, each superpeer  $SP_r$  that receives the clusters  $\{C_s\}$  of another superpeer  $SP_s$  for the first time sends back a *verification* message to  $SP_s$  through the path. The verification message contains the identifiers of the  $SP_s$  and  $SP_r$  superpeers. Each time an intermediate superpeer  $SP_i$  in the path receives a verification message, it adds  $SP_r$  to the corresponding list  $L_s$  and propagates the message back to  $SP_s$ . For example, consider Fig. 7a. The figure depicts (part of) a superpeer network and lets us assume when  $SP_s$  propagates its clusters, the clusters are sent to  $SP_r$  through  $SP_B$ . After  $SP_r$  receives the cluster descriptions from  $SP_s$ , it sends back to  $SP_B$  a verification message, informing  $SP_B$  that the clusters of  $SP_s$  are searchable from  $SP_r$  through  $SP_B$ .

This extra information is useful during query routing, as it avoids redundant forwarding of queries in the network and can be efficiently managed using centralized techniques, like hash tables. The larger the number and the size of cycles, the higher the gain of saved messages. If superpeer  $SP_i$  receives a query for  $SP_s$ 's clusters, the query is forwarded only if the querying superpeer belongs to the list  $L_s$ , otherwise the query is not forwarded further. For

2. Merging bounding boxes of two clusters into a new cluster is a well-studied problem from multidimensional indexing methods. Optimization of the clustering algorithm is out of the scope of this paper; therefore, in our experiments we apply k-Means on the MBRs centers.






Fig. 8. Clusters domination examples. (a)  $C_1 <_U C_2$ . (b)  $C_1 <->_U C_2$ . (c)  $C_1 \geq_U C_2$ .

example, consider a subspace skyline query (Fig. 7b) at  $SP_r$  that has to retrieve data from superpeer  $SP_A$  and  $SP_s$ . Then,  $SP_r$  propagates the query to  $SP_A$  in order to retrieve the necessary data points, and also to  $SP_B$  in order to retrieve the points from  $SP_s$ . Superpeer  $SP_A$  based on its routing information would also forward the query to  $SP_B$  in order to retrieve points from  $SP_s$ . As depicted in Fig. 7, the list  $L_s$  of  $SP_A$  does not contain  $SP_r$ ; therefore,  $SP_A$  knows that  $SP_r$  has contacted  $SP_s$  through another path.

The advantage of using this technique is more efficient query routing, avoiding duplicate messages when the superpeer topology contains cycles. The disadvantage is the maintenance cost, imposed in the case of superpeer failures, for updating the extra (routing) information. In our setup, superpeer failures are infrequent, as superpeers are mainly stable; therefore, our technique is acceptable and provides important gains in performance.

#### 7.4 Query Routing Based on Clusters

The aim of the routing indexes is to identify a set of clusters (and the corresponding superpeers) that may enclose points that belong to the subspace skyline result set. Superpeers are only aware of the cluster descriptions of other superpeers. Therefore, it is necessary to process this information at each superpeer and to retrieve the clusters that may contain nondominated points. In the following, we propose an algorithm (Section 7.4.2) that takes as input the cluster information (instead of points), prunes clusters that cannot provide any results, and returns only those clusters that may contain nondominated points. By knowing only the cluster information, it is feasible to prune parts of other clusters by setting a threshold. Therefore, we assume that a threshold  $t(C_i)$  is attached to each cluster  $C_i$ . In the following, we first discuss (Section 7.4.1) the different cluster relationships and the conditions that may lead to updating a cluster's threshold.

##### 7.4.1 Cluster Dominance Relationships

Based on the employed cluster representations, we can straightforwardly derive dominance relationships between clusters. Given two clusters  $C_1(l_1, u_1)$  and  $C_2(l_2, u_2)$ , we define the following dominance relationships (also depicted in Fig. 8):

**Definition 7.2.** Cluster  $C_1$  dominates  $C_2$  in  $U$ , denoted as  $C_1 <_U C_2$ , if  $l_1$  dominates  $l_2$ .

If  $C_1$  dominates  $C_2$ , this means that there exists at least one point in  $C_1$  that dominates all possible points in  $C_2$ .

Cluster  $C_2$  cannot contribute to the subspace skyline result and can be safely pruned during query processing.

**Definition 7.3.** Clusters  $C_1$  and  $C_2$  are incomparable in  $U$ , denoted as  $C_1 <->_U C_2$ , if  $l_1$  does not dominate  $u_2$  and  $l_2$  does not dominate  $u_1$ .

In this case, no point of  $C_1$  can be dominated nor dominate any point of  $C_2$ , thus both clusters have to be examined during query processing.

**Definition 7.4.** Cluster  $C_1$  partially dominates  $C_2$  in  $U$ , denoted as  $C_1 \geq_U C_2$ , if  $l_1$  dominates  $u_2$  but  $u_1$  does not dominate  $l_2$ .

Therefore, if  $C_1$  partially dominates  $C_2$ , some points of  $C_1$  may dominate some points of  $C_2$ . In this case, we have to examine both clusters in order to retrieve the entire subspace skyline set. However, we can set the threshold of  $C_2$  corresponding to the intersection of the pruning area of  $C_1$  with  $C_2$ .

**Observation 7.5.** Given two clusters  $C_1, C_2$  where  $C_1 \geq_U C_2$ , if  $u_1 <_U u_2$  then the threshold  $t(C_2)$  can be set as  $\text{dist}_U(u_1, C_2)$ .

In the case where  $C_1$  partially dominates  $C_2$  in  $U$ , we can adjust the threshold of  $C_2$ , if  $u_1$  dominates  $u_2$ . Algorithm 4 describes the procedure of updating the thresholds. Each time a threshold is updated, we keep the minimum value of the current threshold and the previously defined threshold by another dominance relationship.

##### Algorithm 4. Threshold update ( $C_n, \{C_i\}$ )

```

1: Input:  $C_n$ : cluster
    $C$ : list of clusters  $\{C_i\}$ 
2: Output: updated thresholds
3: for ( $\forall C_i \in SKY_U : C_n \geq_U C_i$ ) do
4:    $t \leftarrow \max_{j \in U} (u_n[j] - l_i[j])$ 
5:    $t_i \leftarrow \min(t_i, t)$ 
6: end for
7: for ( $\forall C_i \in SKY_U : C_i \geq_U C_n$ ) do
8:    $t \leftarrow \max_{j \in U} (u_i[j] - l_n[j])$ 
9:    $t_n \leftarrow \min(t_n, t)$ 
10: end for
```

The dominance relationships between clusters are essential for our routing algorithm for either pruning entire clusters or restricting the number of points in a cluster that should be accessed, as will be shown shortly.

##### 7.4.2 Superpeer Query Routing

During query processing, each superpeer that receives a subspace skyline query: 1) uses the query routing algorithm (Algorithm 5) to identify local and nonlocal candidate clusters, 2) propagates the query to the neighboring superpeers responsible for the nonlocal clusters, and 3) processes the query locally for the local candidate clusters using Algorithm 3. In the following, we describe the query routing process in detail.

##### Algorithm 5. Superpeer query routing

```

1: Input:  $U$ : query dimensions
    $RC$ : list of routing clusters  $RC_i$ 
2: Output:  $C_{SKY}$ : list of nondominated clusters
```




Fig. 9. Clusters of superpeers  $SP_1$  and  $SP_2$ .

```

3: Sort  $RC$  in ascending order according to the value
    $\min_{j \in U}(rl_i[j])$  for each  $RC_i$ 
4:  $C_{SKY} \leftarrow \emptyset$ 
5:  $t(RC_i) \leftarrow \text{MAX\_INT}, \forall RC_i \in RC$ 
6: while ( $RC \neq \emptyset$ ) do
7:    $RC_m \leftarrow RC.\text{pop}()$ 
8:    $RC \leftarrow RC - \{RC_m\}, RC_m \prec_U RC_i$ 
9:    $\text{cursor} \leftarrow \text{search}([m * c, m * c + f(ru_m, RC_m)])$ 
10:   $C_n \leftarrow \text{cursor.pop}()$ 
11:  while ( $(f(l_n, RC_m) < t(RC_m))$  and  $(C_n \neq \text{null})$ ) do
12:    if ( $\nexists C_i \in C_{SKY} : C_i \prec_U C_n$ ) and
        ( $\nexists RC_i \in RC : RC_i \prec_U C_n$ ) then
13:       $C_{SKY} \leftarrow C_{SKY} - \{C_i\}, C_n \prec_U C_i$ 
14:       $RC \leftarrow RC - \{RC_i\}, C_n \prec_U RC_i$ 
15:       $\text{updateThreshold}(C_n, C_{SKY})$ 
16:       $C_{SKY} \leftarrow C_{SKY} \cup C_n$ 
17:       $\text{updateThreshold}(C_n, RC)$ 
18:    end if
19:     $C_n \leftarrow \text{cursor.pop}()$ 
20:  end while
21: end while
22: return  $C_{SKY}$ 

```

First, let us consider a simple example depicted in Fig. 9. Both superpeers  $SP_1$  and  $SP_2$  eventually store the cluster descriptions of all four clusters. For better readability, we depict only the clusters and not the routing clusters. Let us assume that  $SP_1$  and  $SP_2$  store incomparable clusters (namely  $C_1$  and  $C_4$ ), but also some dominated or partially dominated clusters by other superpeers. If  $SP_1$  receives a 2d skyline query,  $SP_1$  concludes based on the routing clusters that the query has to be propagated to  $SP_2$ , because of the clusters  $C_2$  and  $C_4$ . Also  $SP_1$  starts a local skyline search only for the cluster  $C_1$ , while cluster  $C_3$  is discarded since it is dominated by cluster  $C_4$  of  $SP_2$ . When superpeer  $SP_2$  receives the query, it first examines the routing information, and sets a threshold for cluster  $C_2$ . The threshold is set based on cluster  $C_1$  of  $SP_1$  and corresponds to the pruning area of point  $u_1$ . Thereafter, a skyline search on the clusters  $C_2$  and  $C_4$  is processed. Notice that the threshold for cluster  $C_2$  cannot be refined based on cluster  $C_4$ , even though  $C_2$  is partially dominated by  $C_4$ , since point  $u_4$  does not dominate point  $u_2$ . In this way, each superpeer selectively processes only those clusters, that can affect the result set, using appropriate threshold values.

Algorithm 5 describes the proposed routing algorithm. As in iSUBSKY method, the routing clusters  $RC_i = [rl_i, ru_i]$  are examined in ascending order of the  $\min_{j \in U}(rl_i[j])$  values

(Lines 3-7). During processing, the clusters (Line 13) and routing clusters (Line 14) that are dominated by the currently examined cluster  $C_n$  are discarded. For clusters and routing clusters that are partially dominated by or partially dominate  $C_n$ , we compute new threshold values (Lines 15 and 17, respectively). At the end of Algorithm 5,  $C_{SKY}$  contains all nondominated clusters that should be processed.

A cluster  $C_i \in C_{SKY}$  may be a local cluster and then we add  $C_i$  to the list  $C$  ( $C = C \cup C_i$ ) with clusters that should be processed locally and set the initial threshold of  $C_i$  based on Algorithm 4. Otherwise, the subspace skyline query has to be posed to the corresponding superpeer  $SP(C_i)$ ; therefore, we add  $SP(C_i)$  to the list  $SP$  ( $SP = SP \cup SP(C_i)$ ) of superpeers that should be contacted. Then, the superpeer propagates the query to all  $SP_i \in SP$  and processes the query locally on the list  $C$  by using Algorithm 3.

Notice that in contrast to SKYPEER, we do not need to propagate any threshold among superpeers during query processing. Instead, the threshold is refined at each superpeer based on the routing information, i.e., the cluster descriptions. Similar to SKYPEER, the intermediate results have to be merged as described in Section 6.2, which leads to two different variations of SKYPEER+. In the case of fixed merging at  $P_{init}$ , we have the Routing-based Fixed Merging (RFM) variation, while progressive merging leads to Routing-based Progressive Merging (RPM) variation.

## 8 MAINTENANCE AND COST ANALYSIS

In this section, we present how maintenance is performed with respect to data updates as well as peer joins and failures. Then, we provide a comparative performance analysis for SKYPEER with respect to other existing approaches [5], [6].

### 8.1 Maintenance

In a dynamic P2P network, it is necessary to provide support both for data changes and peer churn. The most common maintenance operations are data insertions and deletions (notice that a data update can be treated as a deletion followed by an insertion). Data insertion at a peer may change the peer's extended skyline; therefore, this leads to an update to the responsible superpeer. In some cases, the insertion causes also the superpeer's extended skyline to change. Data deletions at a peer may change the peer's extended skyline. In those cases, the change needs to be sent to the superpeer, and the superpeer needs to recompute its extended skyline by requesting its peers to send their extended skyline sets.

In the case of SKYPEER, changes at superpeer level do not need to be propagated further, whereas, in the case of SKYPEER+, updates are propagated to the rest of the superpeer network, only if they change the superpeer's MBRs. Therefore, assuming that the data distribution does not change, only a small percentage of data insertions and deletions actually cause communication at superpeer level.

When a peer joins the P2P network by connecting to a superpeer, its extended skyline is transferred to the superpeer and then the same procedure as in data insertions takes place. Peer failures are more complicated, because they resemble data deletion, which means that the responsible superpeer needs to recompute its extended skyline.

TABLE 2  
Performance Analysis

	SKYPEER	Structured P2P
Construction	$O( ext-SKY(S_i) N_P)$	$O( S_i N_P \log N_P)$
Search	$O(N_{sp})$	$O((1+2d(1-\frac{1}{\sqrt{N_P}}))\log N_P)$
Updates	$O(1)$	$O(\log N_P)$

## 8.2 Performance Analysis

In the following, we provide a performance analysis of the SKYPEER framework and an analytical cost comparison with other existing approaches [5], [6] for skyline query processing in P2P networks. We emphasize that the aim is not to present a direct comparison, which is not necessarily meaningful as they rely on different underlying architectures and assumptions, but rather to serve as a comparative performance analysis. Our evaluation uses three important factors that comprise the total cost, namely construction cost, search cost, and update cost.

SKYPEER takes advantage of the two-level architecture of the superpeer system. Therefore, SKYPEER adopts a pre-computation phase at superpeer level, which is similar to materializing results. Other competitor approaches [5], [6] proposed for skyline computation in peer-to-peer systems assume a structured P2P overlay, with distinguishing feature that communication between any two peers is  $O(\log N_P)$ .

In terms of construction cost, SKYPEER requires that each peer sends its extended skyline points to the corresponding superpeer. Therefore, the preprocessing cost of SKYPEER in terms of network traffic, is  $O(|ext-SKY(S_i)| * N_P)$ , where  $|ext-SKY(S_i)|$  represents the cardinality of the local extended skyline of a peer's data set  $S_i$ . On the other hand, structured P2P overlays require for each peer's data point to determine the appropriate peer that indexes the point, which leads to a construction cost of  $O(|S_i| * N_P * \log N_P)$ . The cardinality of the local extended skyline set ( $|ext-SKY(S_i)|$ ) is smaller than the cardinality of the data set ( $|S_i|$ ); therefore, the construction cost of SKYPEER is smaller and it incurs only on local peer to superpeer links.

Regarding search, Wang et al. [5] report an average number of steps equal to:  $O((1 + 2d(1 - \frac{1}{\sqrt{N_P}}))\log N_P)$  for uniform query load. SKYPEER contacts all superpeers during query processing leading to  $O(N_{sp})$  steps (worst case scenario). Therefore, the SKYPEER search cost depends on the number of superpeers ( $N_{sp}$ ), while the cost of [5] depends on the logarithm of the number of peers ( $N_P$ ) in the network. As also indicated by our experiments, these numbers may differ in as much as 1-2 orders of magnitude.

Data insertions or updates have a cost of  $O(\log N_P)$  in structured P2P overlay networks, as in [5]. In SKYPEER, superpeers are informed about data insertions, only if the peer's extended skyline is modified. In those cases, the cost incurred is  $O(1)$  in terms of messages, since the required communication is between a superpeer and its peer.

Concluding, the suitability of each approach clearly depends on the application scenario, as well as on the requirements and characteristics of the P2P system. We summarize the results of our analysis in Table 2.

## 9 EXPERIMENTAL EVALUATION

We evaluate the performance of our algorithms, implemented in Java, using simulations that ran on 3 GHz Pentium IV PCs and locally stored data. In order to be able to test the algorithms with realistic network sizes, we ran multiple peer instances on the same machine and simulated the network interconnection. The P2P network topology used in the experiments consists of  $N_{sp}$  interconnected superpeers in a random graph topology. We used the GT-ITM topology generator<sup>3</sup> to create well-connected random graphs of  $N_{sp}$  peers with a user-specified average connectivity ( $DEG_{sp}$ ). In our experiments, we vary the network size ( $N_p$ ) from 4,000 to 80,000 peers, while the number of superpeers is  $N_{sp} = 5\% \times N_p$  (for  $N_p \geq 20,000$  we used  $N_{sp} = 1\% \times N_p$ ).

We used two different data sets: uniform and clustered. The uniform data set includes random points in a unit space. For the clustered data set, each superpeer picks randomly a point and thereafter  $k = 5$  cluster centroids are generated that follow a Gaussian distribution on each axis with variance 0.025, and a mean equal to the corresponding coordinate of the random point. Then, all associated peers obtain points, the coordinates of which follow a Gaussian distribution on each axis with variance 0.005 around the cluster centroids of the superpeer. Given a query dimensionality, all subspaces have uniform probability to be requested. We generate 100 queries, having a randomly selected superpeer initiator for each query, and we present the average values of our results.

### 9.1 Scalability Study for Uniform Data

In the first series of experiments, we examine the proposed method's scaling features with regards to data set dimensionality. Unless mentioned explicitly, the default values are:  $d = 8$ ,  $k = 3$ ,  $DEG_{sp} = 4$ ,  $N_p = 4,000$ , each peer holds 250 uniformly distributed data points; thus, the cardinality of the data set is  $n = 1M$  data points.

Fig. 10a depicts the performance of SKYPEER in terms of computational time, neglecting network delays. The refined threshold variants (RT\*M) are more costly than the fixed threshold variants (FT\*M), respectively; however, they are still more efficient than the naive one, because of threshold usage. Further, progressive merging (\*TPM) is faster than fixed merging (\*TFM) at the initiator, because the fixed merging at  $P_{init}$  is very costly due to the high number of elements, in contrast to the total merging cost incurring at intermediate superpeers. The plot in Fig. 10b illustrates the total response time, which depends on the size of transmitted data, taking into account the network delay. We assume a modest 4 KB/sec as the network transfer bandwidth on each connection. The four variants of SKYPEER constantly outperform the naive algorithm. Progressive merging reduces communication costs, as several candidate skyline points are pruned out earlier. In Fig. 10c, the volume of the messages (in KB) is presented for various data set dimensionality values, and for query dimensionality 2 and 3. Since the volume of the messages exchanged to retrieve the skyline result depends mainly on

3. Available at: <http://www.cc.gatech.edu/projects/gtitm/>.

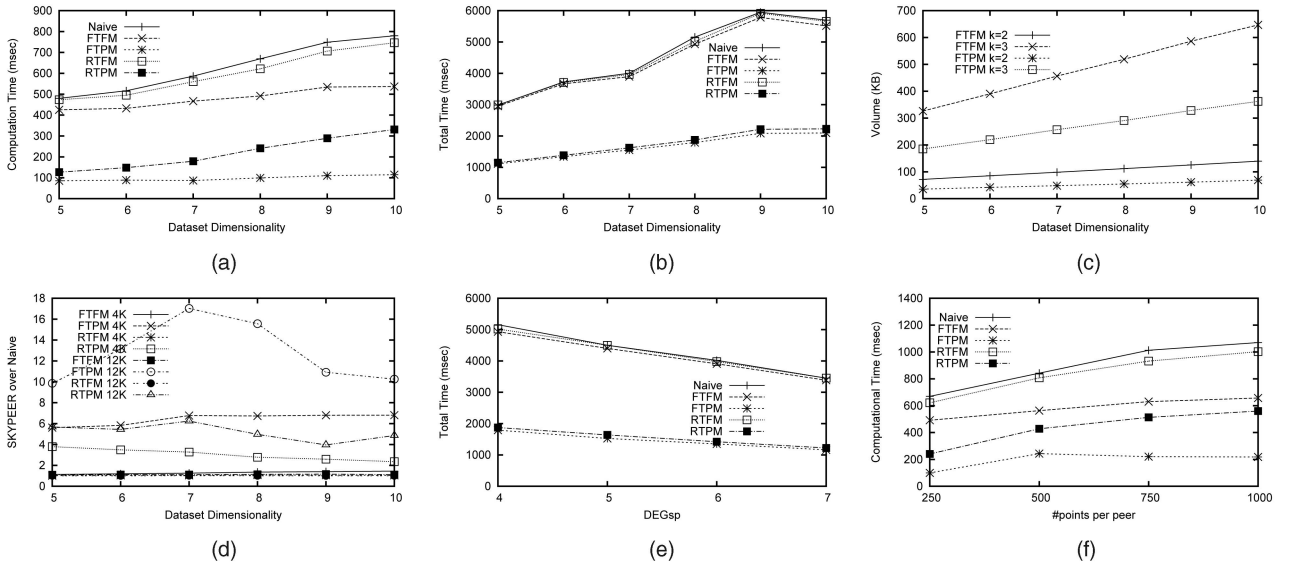


Fig. 10. Comparison of SKYPEER variants for uniform data distribution. (a) Computational time. (b) Total time. (c) Volume of transferred messages. (d) SKYPEER's relative performance. (e) Total time for varying  $DEG_{sp}$ . (f) Computational time for varying #points/peer.

the merging strategy, we present the results of FTFM and FTFM as representative ones. Obviously progressive merging reduces the volume of transferred data.

In order to study SKYPEER's behavior for larger networks, we increased the network size from 4,000 to 12,000 peers. In Fig. 10d, we present a performance comparison between SKYPEER and the naive algorithm for different network sizes. The vertical axis represents the relative performance of SKYPEER as compared to the naive approach. It is clear that SKYPEER always outperforms naive, and for large networks (12,000 peers) the FTFM variant is 17 times faster than naive. We also study how different  $DEG_{sp}$  values affect the performance of SKYPEER. In Fig. 10e, we experiment with a network of 4,000 peers with  $DEG_{sp}$  varying from 4 to 7. We conclude that even though the computational time is not affected by  $DEG_{sp}$ , the total time is reduced when  $DEG_{sp}$  increases. This is because higher  $DEG_{sp}$  values, result in

smaller routing paths, hence smaller network transfer costs. Finally, in Fig. 10f, we increased the number of points per peer ( $n/N_p$ ) from 250 to 1,000. We notice that the progressive merging variants clearly outperform the fixed merging ones, as the number of points per peer increases.

## 9.2 Scalability Study for Clustered Data

In the next experiments, we evaluate the routing ability of SKYPEER<sup>+</sup> using a clustered data set distributed over  $N_p = 4,000$  peers. Each peer holds 250 points. In Fig. 11a, we depict the response time for the best variants, namely refine threshold (RT\*M) and routing-based (R\*M). We notice that the response time for the progressive merging methods is slightly higher in the case of RTPM than for RPM. However, the fixed merging approach RFM performs much better than RTFM. This is mainly caused by the very small number of

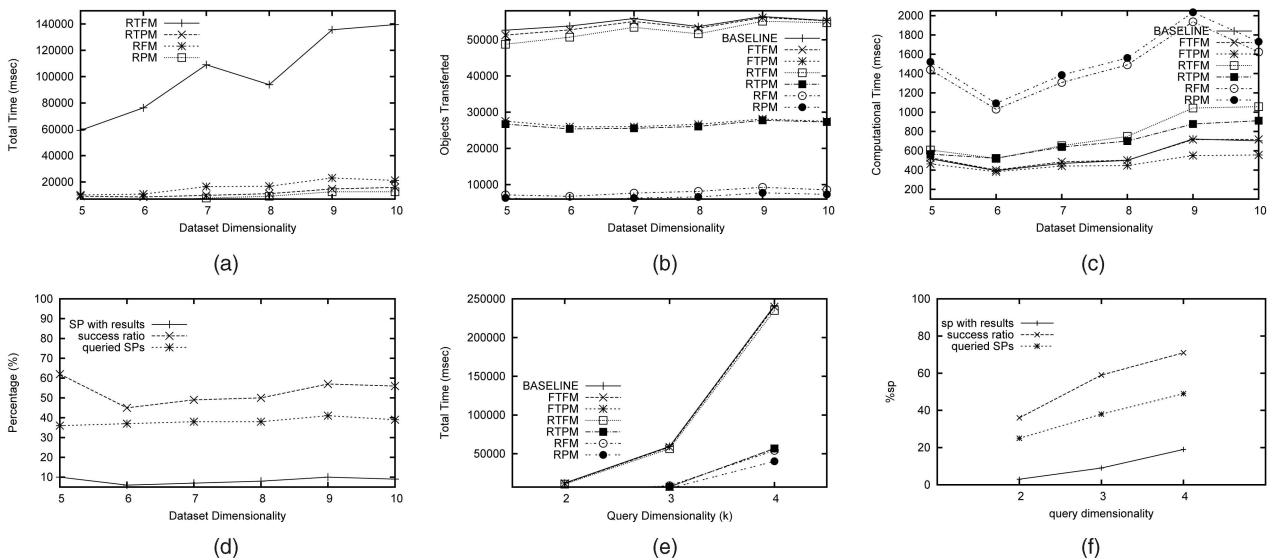


Fig. 11. Comparison of SKYPEER<sup>+</sup> with SKYPEER variants for clustered data distribution. (a) and (e) Total time. (b) Transferred objects. (c) Computational time. (d) and (f) Success ratio.

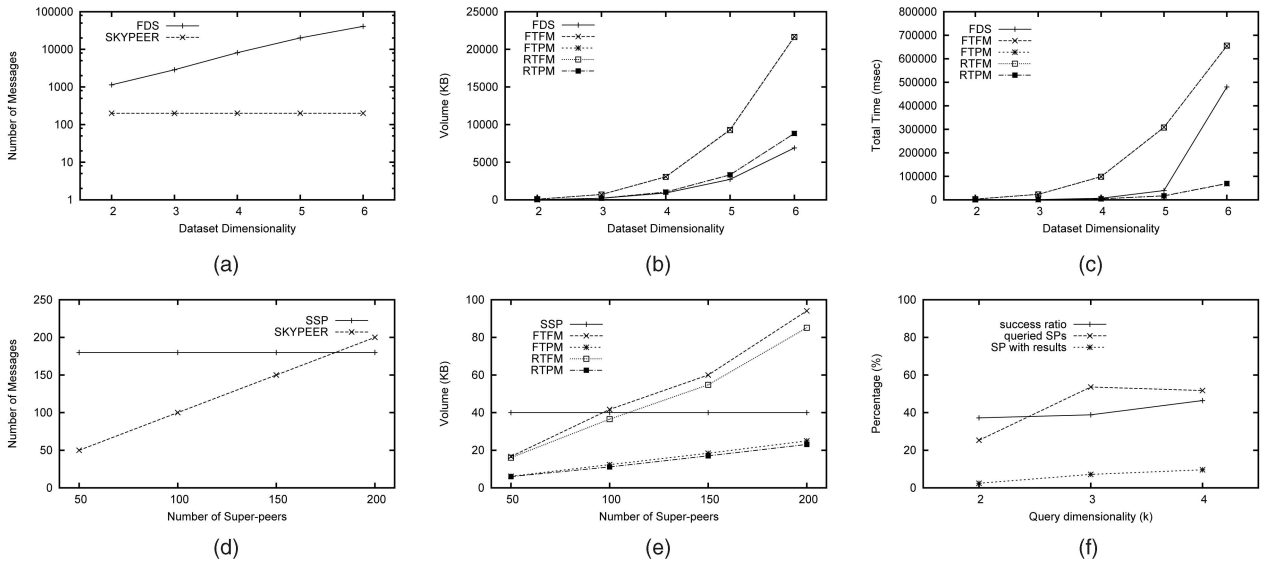


Fig. 12. Comparison of SKYPEER with other algorithms (FDS [8], SSP [5]) and experiment with real data. (a) Number of Messages (versus FDS). (b) Volume of transf. messages (versus FDS). (c) Total time (versus FDS). (d) Number of Messages (versus SSP). (e) Volume of transf. messages (versus SSP). (f) Success ratio for real data.

objects that are transferred (Fig. 11b). For the routing-based approach, the threshold is refined for each superpeer based on the routing indexes that summarize the information available on the whole network. In Fig. 11b, we notice that the number of transferred objects is very low, especially in the case of RFM that has only few more objects more than RPM but much less (order of magnitude) than the other fixed merging approaches. On the other hand, as depicted in Fig. 11c, the fixed threshold variants need less computational time than the refined variants and even less than the routing variants. This additional cost is caused by the processing cost induced by the routing indexes; however, it is negligible compared to the total response time. Notice that the merging phase does not influence significantly the computational time, since the merging time is low.

Fig. 11d depicts: 1) the percentage of superpeers that store at least one point that belongs to the overall subspace skyline set, 2) the percentage of superpeers that are contacted when the subspace skyline query is routed, and 3) the percentage of the contacted superpeers that could not have been avoided to be queried, either because they store some subspace skyline points or because they have to propagate the query to some superpeer that stores some skyline points. In Fig. 11e, we vary the query dimensionality from 2 to 4 and we measure the response time for all variations. For any query dimensionality, SKYPEER<sup>+</sup> performs better than the SKYPEER variants, and the gain increases with the query dimensionality. Analogously to Fig. 11d, 11f shows for varying query dimensionality that the success ratio increases up to 70 percent, while the queried superpeers are less than 50 percent.

### 9.3 Comparison to Existing Algorithms

In the following, we provide a comparative study against two approaches (FDS algorithm [8] and SSP [5]) for distributed skyline computation, even though they assume different distributed architectures. First, we compare SKYPEER

against FDS,<sup>4</sup> which assumes a set of servers and the absence of a network overlay. In our scenario, each server is a superpeer and even if SKYPEER assumes an overlay network, for the FDS approach, we assume that each server communicates directly with any other server. We emphasize that this experimental setup favors FDS, as any two servers communicate with minimum cost without the need for efficient query routing. Figs. 12a, 12b, and 12c show the number of search messages, their volume, and the response time, respectively, for SKYPEER and FDS for increasing dimensionality. We employ a network of 2,000 peers (with  $DEG_p = 10$  in the case of SKYPEER),  $n = 1M$ , we use the uniform data set and we evaluate full space skyline queries. Notice that SKYPEER requires fewer messages than FDS (Fig. 12a), because FDS uses several communication phases to compute the result. In terms of transferred volume (Fig. 12b), FDS is cheaper than SKYPEER, since FDS is designed to minimize the number of transferred objects. However, notice that the \*TPM variants of SKYPEER achieve comparable results to FDS, due to the efficiency of in-network progressive merging. Furthermore, the transferred volume of FDS would increase rapidly, if FDS would follow the existing paths for transferring data, instead of direct communication between superpeers. Moreover, in Fig. 12c, the \*TPM variants of SKYPEER outperform FDS in terms of total response time, when the dimensionality increases. This is because the total time includes the processing costs, which in the case of FDS are much higher than \*TPM, especially for higher dimensions.

In the next experiment, we compare against the SSP [5] algorithm, that relies on a structured P2P overlay network. Structured P2P overlays relocate data on the peers, which is beneficial for query processing. On the other hand, super-peer networks gain in query processing performance due to degree of centralization introduced by the superpeers. In

4. Notice that the algorithm mentioned as ALS in [8] is the baseline variant of SKYPEER.

Figs. 12d and 12e, we compare our approach with SSP [5] by varying the degree of centralization of SKYPEER. We use a network of 2K peers,  $n = 1.6M$  and a two-dimensional uniform data set. We gradually increase the number of superpeers ( $N_{sp}$ ) employed by SKYPEER. In Fig. 12d, the number of search messages required by SSP is constant as it does not depend on  $N_{sp}$ . However, depending to the number of superpeers employed, SKYPEER performs better or worse than SSP. For small numbers of superpeers, SKYPEER needs fewer messages than SSP, but when the number of superpeers increases, the performance of SKYPEER degrades. Nevertheless, when studying the transferred volume (Fig. 12e), the \*TPM variants of SKYPEER are always better than SSP. Again, this is because the in-network progressive merging discards several results before they reach the querying peer.

In summary, the conclusions of the comparative study are:

- SKYPEER is more efficient than FDS in terms of number of required messages. Furthermore, SKYPEER scales better than FDS with dimensionality. FDS transfers marginally fewer data; however, the \*TPM variants of SKYPEER have similar performance.
- The \*TPM variants of SKYPEER achieve smaller total response time than FDS and this gain increases as dimensionality increases.
- Compared to SSP, SKYPEER requires fewer messages for query processing, when the number of superpeers is smaller than 9 percent of the number of peers. Moreover, the \*TPM variants of SKYPEER induce smaller volume of transferred objects, for all tested setups.

## 9.4 Experimental Evaluation with Real Data

We also evaluate SKYPEER<sup>+</sup> using real-life data. We crawled data (from [www.zillow.com](http://www.zillow.com)) containing information about real estate all over the United States. We obtained a six-dimensional data set containing 1M entries. The data set contains six attributes, namely number of bathrooms, number of bedrooms, living area, price, year built, and lot area. We distributed the data set randomly to 4,000 peers and we use 200 superpeers. Fig. 12f depicts the results for varying query dimensionality  $k$ . The success ratio increases up to 46.5 percent, and the queried superpeers are less than 50 percent. These results verify the performance of SKYPEER<sup>+</sup> also for real-life data.

## 10 CONCLUSIONS

In this paper, we addressed the problem of efficient skyline query routing over a P2P network. We presented a threshold-based algorithm, called SKYPEER, which forwards the skyline query requests among peers, in such a way that the amount of data to be transferred is significantly reduced. Thereafter, we proposed SKYPEER<sup>+</sup>, an advanced routing algorithm, which improves performance in the case of clustered data distribution at superpeers. SKYPEER<sup>+</sup> employs an appropriate indexing technique, in order to maintain routing information at superpeer level. Finally, in our experimental evaluation, we demonstrate the efficiency of our approach both in terms of computational and communication costs.

## ACKNOWLEDGMENTS

This work was carried out during the tenure of an ERCIM “Alain Bensoussan” Fellowship Programme that supported Akrivi Vlachou and Christos Doulieridis.

## REFERENCES

- [1] S. Börzsönyi, D. Kossmann, and K. Stocker, “The Skyline Operator,” *Proc. Int’l Conf. Data Eng. (ICDE)*, pp. 421-430, 2001.
- [2] W.-T. Balke, U. Gunzer, and J. Zheng, “Efficient Distributed Skylining for Web Information Systems,” *Proc. Int’l Conf. Extending Database Technology (EDBT)*, pp. 256-273, 2004.
- [3] B. Cui, H. Lu, Q. Xu, L. Chen, Y. Dai, and Y. Zhou, “Parallel Distributed Processing of Constrained Skyline Queries by Filtering,” *Proc. Int’l Conf. Data Eng. (ICDE)*, pp. 546-555, 2008.
- [4] A. Vlachou, C. Doulieridis, Y. Kotidis, and M. Vazirgiannis, “SKYPEER: Efficient Subspace Skyline Computation over Distributed Data,” *Proc. Int’l Conf. Data Eng. (ICDE)*, pp. 416-425, 2007.
- [5] S. Wang, B.C. Ooi, A.K.H. Tung, and L. Xu, “Efficient Skyline Query Processing on Peer-to-Peer Networks,” *Proc. Int’l Conf. Data Eng. (ICDE)*, pp. 1126-1135, 2007.
- [6] P. Wu, C. Zhang, Y. Feng, B.Y. Zhao, D. Agrawal, and A.E. Abbadi, “Parallelizing Skyline Queries for Scalable Distribution,” *Proc. Int’l Conf. Extending Database Technology (EDBT)*, pp. 112-130, 2006.
- [7] L. Zhu, S. Zhou, and J. Guan, “Efficient Skyline Retrieval on Peer-to-Peer Networks,” *Future Generation Comm. and Networking*, vol. 1, pp. 309-314, 2007.
- [8] L. Zhu, Y. Tao, and S. Zhou, “Distributed Skyline Retrieval with Low Bandwidth Consumption,” *IEEE Trans. Knowledge and Data Eng.*, vol. 21, no. 3, pp. 384-400, Mar. 2009.
- [9] C. Doulieridis, A. Vlachou, Y. Kotidis, and M. Vazirgiannis, “Peer-to-Peer Similarity Search in Metric Spaces,” *Proc. Int’l Conf. Very Large Data Bases (VLDB)*, pp. 986-997, 2007.
- [10] C. Doulieridis, K. Nøravåg, and M. Vazirgiannis, “DESENT: Decentralized and Distributed Semantic Overlay Generation in P2P Networks,” *IEEE J. Selected Areas in Comm.*, vol. 25, no. 1, pp. 25-34, Jan. 2007.
- [11] H.T. Kung, F. Luccio, and F.P. Preparata, “On Finding the Maxima of a Set of Vectors,” *J. ACM*, vol. 22, no. 4, pp. 469-476, 1975.
- [12] D. Kossmann, F. Ramsak, and S. Rost, “Shooting Stars in the Sky: An Online Algorithm for Skyline Queries,” *Proc. Int’l Conf. Very Large Data Bases (VLDB)*, pp. 275-286, 2002.
- [13] J. Chomicki, P. Godfrey, J. Gryz, and D. Liang, “Skyline with Pre-Sorting,” *Proc. Int’l Conf. Data Eng. (ICDE)*, pp. 717-719, 2003.
- [14] D. Papadias, Y. Tao, G. Fu, and B. Seeger, “Progressive Skyline Computation in Database Systems,” *ACM Trans. Database Systems*, vol. 30, no. 1, pp. 41-82, 2005.
- [15] P. Godfrey, R. Shipley, and J. Gryz, “Maximal Vector Computation in Large Data Sets,” *Proc. Int’l Conf. Very Large Data Bases (VLDB)*, pp. 229-240, 2005.
- [16] I. Bartolini, P. Ciaccia, and M. Patella, “SaLSa: Computing the Skyline without Scanning the Whole Sky,” *Proc. ACM Int’l Conf. Information and Knowledge Management (CIKM)*, pp. 405-414, 2006.
- [17] J. Pei, W. Jin, M. Ester, and Y. Tao, “Catching the Best Views of Skyline: A Semantic Approach Based on Decisive Subspaces,” *Proc. Int’l Conf. Very Large Data Bases (VLDB)*, pp. 253-264, 2005.
- [18] Y. Yuan, X. Lin, Q. Liu, W. Wang, J. Yu, and Q. Zhang, “Efficient Computation of the Skyline Cube,” *Proc. Int’l Conf. Very Large Data Bases (VLDB)*, pp. 241-252, 2005.
- [19] Y. Tao, X. Xiao, and J. Pei, “SUBSKY: Efficient Computation of Skylines in Subspaces,” *Proc. Int’l Conf. Data Eng. (ICDE)*, 2006.
- [20] Y. Tao, X. Xiao, and J. Pei, “Efficient Skyline and Top-K Retrieval in Subspaces,” *IEEE Trans. Knowledge and Data Eng.*, vol. 19, no. 8, pp. 1072-1088, Aug. 2007.
- [21] K. Hose, C. Lemke, and K.-U. Sattler, “Processing Relaxed Skylines in PDMS Using Distributed Data Summaries,” *Proc. ACM Int’l Conf. Information and Knowledge Management (CIKM)*, pp. 425-434, 2006.
- [22] Z. Huang, C. Jensen, H. Lu, and B.-C. Ooi, “Skyline Queries against Mobile Lightweight Devices in MANETs,” *Proc. Int’l Conf. Data Eng. (ICDE)*, 2006.
- [23] S. Wang, Q.H. Vu, B.C. Ooi, A.K.H. Tung, and L. Xu, “Skyframe: A Framework for Skyline Query Processing in Peer-to-Peer Systems,” *Very Large Data Bases J.*, vol. 18, no. 1, pp. 345-362, 2009.

- [24] K. Fotiadou and E. Pitoura, "Bitpeer: Continuous Subspace Skyline Computation with Distributed Bitmap Indexes," *Proc. Int'l Workshop Data Management in Peer-to-Peer Systems (DAMAP)*, 2008.
- [25] A. Crespo and H. Garcia-Molina, "Routing Indices for Peer-to-Peer Systems," *Proc. Int'l Conf. Distributed Computing Systems (ICDCS)*, pp. 23-34, 2002.
- [26] B. Yang and H. Garcia-Molina, "Improving Search in Peer-to-Peer Networks," *Proc. Int'l Conf. Distributed Computing Systems (ICDCS)*, pp. 5-14, 2002.
- [27] H.V. Jagadish, B.C. Ooi, K.-L. Tan, C. Yu, and R. Zhang, "iDistance: An Adaptive B+-Tree Based Indexing Method for Nearest Neighbor Search," *ACM Trans. Database Systems*, vol. 30, no. 2, pp. 364-397, 2005.



**Akrivi Vlachou** received the BSc and MSc degrees from the Department of Computer Science and Telecommunications at the University of Athens in 2001 and 2003, respectively. She received the PhD degree in 2008 from the Athens University of Economics and Business (AUEB). She is currently a postdoctoral researcher at the Norwegian University of Science and Technology (NTNU). She was recently awarded with an ERCIM "Alain Bensoussan" fellowship. Her research interests include query processing and data management in large-scale distributed systems. She is a member of the IEEE.



**Christos Doulkeridis** received the BSc degree in electrical engineering and computer science from the National Technical University of Athens, and the MSc and PhD degrees in information systems from the Department of Informatics at the Athens University of Economics and Business. He is currently a postdoctoral researcher with an ERCIM "Alain Bensoussan" fellowship at the Norwegian University of Science and Technology. His research interests include data management in P2P networks, distributed knowledge discovery, and mobile and context-aware computing. He is a member of the IEEE.



**Yannis Kotidis** received the BSc degree in electrical engineering and computer science from the National Technical University of Athens, and the MSc and PhD degrees in computer science from the University of Maryland. He is an assistant professor in the Department of Informatics at the Athens University of Economics and Business. Between 2000 and 2006, he was a senior technical specialist at the Database Research Department of AT&T Labs-Research in Florham Park, New Jersey. His main research areas include large-scale data management systems, data warehousing, and sensor networks.



**Michalis Vazirgiannis** received the degree in physics (1986) and the MSc degree in robotics (1988), both from the University of Athens, and the MSc degree in knowledge-based systems from Heriot Watt University in Edinburgh, United Kingdom. He received the PhD degree in 1994 from the Department of Informatics at the University of Athens, Greece. He is an associate professor in the Department of Informatics at the Athens University of Economics and Business. Since then, he has been conducting research at GMD-IPSI, Max Planck MPI in Germany, and INRIA/FUTURS in Paris, France. He has been a visiting professor at Ecole Polytechnique in France and Deusto University in Spain. His research interests and work range from interactive multimedia systems to spatiotemporal databases and distributed data/web mining.

► **For more information on this or any other computing topic, please visit our Digital Library at [www.computer.org/publications/dlib](http://www.computer.org/publications/dlib).**